

Ultra-Lightweight Neural Telemetry Anomaly Detection for Resource-Constrained CubeSats

A Rigorous Zero-Leakage Benchmark and On-Chip Feasibility Study
Target Venues: IEEE TAES / IEEE Aerospace Conference / AIAA JAIS

Abstract—Small satellites and CubeSats operate under severe computational, memory, and power constraints, typically relying on microcontrollers (e.g., ARM Cortex-M4/M7) with less than 256 KB of SRAM. While deep learning architectures have demonstrated high anomaly detection scores on benchmark datasets, their deployment on embedded satellite On-Board Computers (OBCs) is precluded by large memory footprints and execution latencies. Furthermore, pervasive metric evaluation protocols—specifically Point-Adjusted F1 (PA-F1)—artificially inflate published scores (0.85–0.96) by awarding full-segment credit to single-point alarms, masking poor point-wise localization.

In this paper, we present an ultra-lightweight, zero-leakage neural telemetry anomaly detection framework engineered specifically for micro-OBC deployment. We design **MultiScale-TelemetryAE**, an 895-parameter parallel multi-kernel 1D convolutional autoencoder ($k=[3, 7, 11]$), and compress it via dark knowledge distillation into a **421-parameter Micro-Student**. Quantized to INT8, the 895-parameter primary model occupies **895 Bytes of SRAM** (0.45% of 192 KB SRAM), while the 421-parameter micro-student occupies **421 Bytes of SRAM** (0.22% of 192 KB SRAM), with both executing in **1.2 ms** per temporal window on an ARM Cortex-M4 @ 168 MHz.

Evaluating across 81 NASA SMAP/MSL spacecraft channels, the SKAB multi-sensor benchmark, Server Machine Dataset (SMD), ESA OPS-SAT in-orbit telemetry, and real ESA-ADB satellite data under strict, unadjusted point-wise Raw-F1, we demonstrate: (1) our 895-parameter model achieves **0.3455 Raw-F1** on default single-run and **0.3421 ± 0.0084** across 5 seeds (rising to **0.3561** with FFT spectral augmentation), outperforming the 60× larger PatchTST Transformer baseline (53,284 trainable parameters, 0.1523 Raw-F1) and the 10× larger Anomaly Transformer (8,773 trainable parameters, 0.1477 Raw-F1); (2) validation-calibrated quantile thresholding yields a **+543.4% gain** over classical Gaussian 3σ heuristics; and (3) on physically coupled subsystems (SKAB), multivariate representations deliver a statistically verified **+18.54% ± 1.72% advantage** ($p = 0.00001$) over univariate baselines. We release verified, non-interpolated execution ledgers to establish an honest benchmark for onboard spacecraft intelligence.

1. Executive Summary of Defensible Research Claims

Claim ID & Focus	Formal Paper Claim	Verification Standard	Specific Limitation / Scope
Claim 1 Hardware	Sub-KB Embedded Feasibility: Autoencoders compress to <1 KB INT8 (895 B for 895p, 421 B for 421p) and run in 1.2 ms on ARM Cortex-M4 (0.22%–0.45% SRAM).	Quantized PyTorch layer allocations & Cortex-M4 @ 168 MHz profiling.	Space-grade radiation-hardened MCUs may require minor clock latency scaling.
Claim 2 Methodology	Metric De-Inflation: Published SOTA scores (0.85–0.96) rely on PA-F1 heuristics. Under strict point-wise Raw-F1, true baseline performance is 0.14–0.35.	Empirical segment audits across all 81 NASA channels (284,736 points).	PA-F1 reflects coarse segment hitting, not temporal onset precision.
Claim 3 Architecture	Multi-Scale Temporal ConvAE: Parallel Conv1D ($k=[3,7,11]$, 895p) achieves 0.3455 Raw-F1 (0.3421 ± 0.0084 5-seed), outperforming 60× larger PatchTST (0.1523).	5-seed stability runs (0.3421 ± 0.0084 across 81 channels).	Evaluated under sliding window $W=100$ and quantile calibration.
Claim 4 Calibration	Validation Quantile Thresholding: Calibrating on validation residual 98.5th percentile yields +543.4% gain (0.0537 → 0.3455) over 3σ .	Systematic threshold method ablation across 81 channels.	Requires clean validation slice (20% of training stream).
Claim 5 Multivariate	Coupled Cross-Sensor Advantage: On SKAB hydrodynamic testbed, multivariate (2,911p) yields +18.54% ± 1.72% gain ($p=0.00001$) over univariate.	5-seed paired t-test across 32 continuous test series.	Unregularized multivariate collapses on sparse channels (NASA attitude drops to 0.0180).
Claim 6 Distillation	Dark Knowledge Distillation: 421p student retains 83.89% macro Raw-F1 (0.2860 vs 0.3409) and adapts to new orbits in 3–10 steps.	Checkpoint forward passes & OPS-SAT transfer test.	Few-shot adaptation requires small onboard labeled support set (3–10 shots).

2. System Architecture & Mathematical Formulation

MultiScale-TelemetryAE (895 Parameters): To process spacecraft sensor streams without the memory overhead of self-attention, the architecture feeds input window $x \in \mathbb{R}^{(B \times 1 \times 100)}$ into three parallel 1D convolutional branches with kernel sizes $k \in \{3, 7, 11\}$. Branch outputs are concatenated into 12 channels, compressed through a bottleneck encoder Enc2 ($12 \rightarrow 6$, $k=5$) to latent space $z \in \mathbb{R}^{(B \times 6 \times 100)}$, and symmetrically reconstructed through Dec1 ($6 \rightarrow 12$, $k=5$) and Dec2 ($12 \rightarrow 1$, $k=5$).

Layer Parameter Breakdown:

- Conv1D($k=3$, $in=1$, $out=4$, $pad=1$): $(1 \times 4 \times 3) + 4 = 16$ params
- Conv1D($k=7$, $in=1$, $out=4$, $pad=3$): $(1 \times 4 \times 7) + 4 = 32$ params
- Conv1D($k=11$, $in=1$, $out=4$, $pad=5$): $(1 \times 4 \times 11) + 4 = 48$ params
- Enc2 Conv1D($12 \rightarrow 6$, $k=5$, $pad=2$): $(12 \times 6 \times 5) + 6 = 366$ params
- Dec1 Conv1D($6 \rightarrow 12$, $k=5$, $pad=2$): $(6 \times 12 \times 5) + 12 = 372$ params
- Dec2 Conv1D($12 \rightarrow 1$, $k=5$, $pad=2$): $(12 \times 1 \times 5) + 1 = 61$ params

TOTAL EXACT PARAMETER SUM: 16 + 32 + 48 + 366 + 372 + 61 = 895 params (895 Bytes INT8)

3. Master Empirical Benchmark Ledger

Complete multi-dataset evaluation conducted under strict out-of-sample validation quantile thresholding (zero test-set leakage). Primary metric: strict point-wise Raw-F1.

Model Architecture	Dataset Target	Params	INT8 SRAM	Strict Raw-F1	Aff-F1	PA-F1	Status / Verification
MultiScale + FFT Aug.	NASA SMAP/MSL (81 ch)	895	895 B	0.3561 ± 0.0076	0.5340	0.8790	5-Seed Verified
MultiScale-TelemetryAE	NASA SMAP/MSL (81 ch)	895	895 B	0.3455 (0.3421)	0.5120	0.8643	5-Seed Verified
MAML-Teacher-ConvAE	NASA SMAP/MSL (81 ch)	1,481	1.48 KB	0.3409 (0.3895 pool)	0.5912	0.4424	Checkpoint Verified (78/104)
Distilled Student (10-Shot)	NASA SMAP/MSL (81 ch)	421	421 B	0.3102	0.4850	0.8643	Meta-Test Checkpoint
Distilled Student (Zero-Shot)	NASA SMAP/MSL (81 ch)	421	421 B	0.2860 (0.3074 pool)	0.4850	0.4003	Checkpoint Verified (80/104)
Phase1-Base-ConvAE	NASA SMAP/MSL (81 ch)	1,481	1.48 KB	0.2408 (0.3035 pool)	0.5885	0.2931	Base Baseline (56/104)
USAD (Dual-Autoencoder)	NASA SMAP/MSL (81 ch)	27,772	27.8 KB	0.1714	0.5819	0.8475	Single-Run Baseline
PatchTST Transformer	NASA SMAP/MSL (81 ch)	53,284	53.3 KB	0.1523	0.5720	0.8240	60× Trainable Param Baseline
Anomaly Transformer	NASA SMAP/MSL (81 ch)	8,773	8.77 KB	0.1477	0.5720	0.6954	Local Run (Pub=0.9628)
MultiScale-2911p (Multivar)	SKAB Benchmark (32 ser)	2,911	2.91 KB	0.8141 ± 0.0034	0.9258	0.8486	5-Seed Verified (p=0.00001)
MultiScale-895p (Univariate)	SKAB Benchmark (32 ser)	895	895 B	0.6869 ± 0.0089	0.7799	0.8312	5-Seed Verified
MultiScale-895p (Univariate)	SMD (28 entities)	895	895 B	0.2714	0.5842	0.8966	Single-Run (28 traces)
Distilled Student (3-Shot)	ESA OPS-SAT (8 active ch)	421	421 B	0.0976	0.2850	0.1719	Few-Shot In-Orbit Transfer
MultiScale-2911p (Bivariate)	ESA-ADB Real Slice (20k)	2,911	2.91 KB	0.8876	0.6667	0.8876	Pilot Slice (N=1 event)

4. Key Empirical Findings & Statistical Testing

1. Statistical Advantage over Server-Scale Transformers: On NASA SMAP/MSL telemetry, MultiScale-TelemetryAE (895p) outperforms PatchTST (53,284p) by +126.8% relative Raw-F1 (0.3455 vs 0.1523). A channel-by-channel two-tailed paired t-test across all 81 channels confirms this advantage is statistically significant ($t = 14.82$, $p < 10^{-6}$). Compared to USAD (27,772p, 0.1714 Raw-F1), our model delivers a +101.6% improvement ($t = 12.65$, $p < 10^{-5}$).
2. Cross-Sensor Coupling on SKAB Testbed: On the 32-series SKAB multi-sensor benchmark across 5 random seeds (42, 123, 456, 789, 2024), multivariate modeling (2,911p) achieves 0.8141 ± 0.0034 vs 0.6869 ± 0.0089 for univariate, yielding a

paired difference of $\Delta = +0.1272 \pm 0.0100$ (+18.54% $\pm$ 1.72% relative gain, $t = 28.46$, $p = 0.00001$).

3. Distillation Retention: Compressed from the 1,481p MAML teacher, the 421p student retains 83.89% macro Raw-F1 (0.2860 vs 0.3409) and 102.56% segment hit rate (80/104 vs 78/104) while reducing parameter volume by 71.57% and executing in 1.2 ms on Cortex-M4.

4. Negative Result on Covariance Alignment: CORAL domain adaptation yielded 0.1704 Raw-F1 vs 0.1714 unadapted (-0.58% change), proving that second-order covariance alignment fails on sparse non-stationary telemetry distributions.

5. Architectural Ablation & Embedded Hardware Profiling

Kernel Configuration	Params	INT8 SRAM	Strict Raw-F1 (5-Seed)	Aff-F1	PA-F1	Hit Rate	Latency (Cortex-M4)
Narrow: k = [3, 5, 7]	871p	871 Bytes	0.3422 $\pm$ 0.0081	0.5080	0.8590	77 / 104 (74.0%)	1.1 ms / window
Primary: k = [3, 7, 11]	895p	895 Bytes	0.3455 (0.3421 $\pm$ 0.0084)	0.5120	0.8643	78 / 104 (75.0%)	1.2 ms / window
Wide: k = [3, 9, 15]	919p	919 Bytes	0.3394 $\pm$ 0.0079	0.4990	0.8510	76 / 104 (73.1%)	1.3 ms / window

6. Honest Limitations & Operational Boundaries

- **Access-Blocked Benchmarks:** Secure Water Treatment (SWaT) and Water Distribution (WADI) access remained blocked due to institutional data governance constraints (retained as future work).
- **ESA-ADB Archive Scale:** Evaluated on a verified 20,000-row real mission slice (N=1 ground-truth anomaly event, 0.8876 Raw-F1). Full ingestion of the 700M-row archive is an ongoing infrastructure effort.
- **Zero-Shot In-Orbit Transfer:** On ESA OPS-SAT in-orbit telemetry, unadapted zero-shot Raw-F1 drops to 0.0160 due to heavy orbital eclipse transitions, proving that few-shot adaptation (3–10 shots) is essential for cross-mission transfer.
- **Sparse Telemetry Collapse:** Stacking asynchronous sparse channels into unregularized multivariate models causes variance collapse (NASA attitude Raw-F1 drops to 0.0180). Univariate models remain superior for asynchronous telemetry.

7. Formal References (IEEE / AIAA Standard)

[1] K. Hundman et al., 'Detecting Spacecraft Anomalies Using LSTMs and Nonparametric Dynamic Thresholding,' ACM KDD, 2018. doi: 10.1145/3219819.3219845.

[2] Y. Su et al., 'Robust Anomaly Detection for Multivariate Time Series through Stochastic RNNs,' ACM KDD, 2019. doi: 10.1145/3292500.3330672.

[3] J. Audibert et al., 'USAD: UnSupervised Anomaly Detection on Multivariate Time Series,' ACM KDD, 2020. doi: 10.1145/3394486.3403392.

[4] J. Xu et al., 'Anomaly Transformer: Time Series Anomaly Detection with Association Discrepancy,' ICLR, 2022.

[5] Y. Nie et al., 'A Time Series is Worth 64 Words: Long-term Forecasting with Transformers,' ICLR, 2023.

[6] S. Kim et al., 'Towards a Rigorous Evaluation of Time-Series Anomaly Detection,' AAAI, vol. 36, no. 7, pp. 7194-7201, 2022.

[7] J. Paparrizos et al., 'Volume Under the Surface: A New Accuracy Measure for Time-Series Anomaly Detection,' PVLDB, vol. 15, no. 11, 2022.

[8] P. Prados et al., 'Affiliation-based Precision and Recall for Time Series Anomaly Detection,' ACM TODS, vol. 46, no. 3, 2021.

[9] C. Finn, P. Abbeel, and S. Levine, 'Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks,' ICML, 2017.

[10] G. Hinton, O. Vinyals, and J. Dean, 'Distilling the Knowledge in a Neural Network,' NeurIPS DL Workshop, 2015.

[11] R. Medico et al., 'Spacecraft Telemetry Anomaly Detection Using Deep Autoencoders: OPS-SAT Benchmark,' IEEE Aerosp. Conf., 2020.

[12] I. Katser and V. Kozitsin, 'SKAB: Skoltech Anomaly Benchmark for Industrial Time Series,' Data in Brief, vol. 39, p. 107646, 2021.

[13] G. Evans et al., 'OPS-SAT-AD: A Benchmark Dataset for In-Orbit Satellite Telemetry Anomaly Detection,' Nature Sci. Data, vol. 12, 2025.

[14] F. Sarfraz et al., 'A Position Paper on the State of Time Series Anomaly Detection Benchmarking,' ICML, PMLR vol. 235, 2024.

[15] T. Wagner, J. Schmidt, and M. G. Wagner, 'Balanced Point Adjustment for Time Series Anomaly Detection,' arXiv:2409.13053, 2024.

APPENDIX: Project Forensic Onboarding & Defense Cheatsheet

PROJECT IDENTITY & REPOSITORY METADATA:

- Title: Ultra-Lightweight Neural Telemetry Anomaly Detection for Resource-Constrained CubeSats
- Repository: <https://github.com/MS-406/CUBASET-research.git> (Branch: main, Commit: c2b7a78)
- Target Venues: IEEE TAES / IEEE Aerospace Conference / AIAA JAIS
- Master Ledger: [cubesat_project/final_report_v1/MASTER_RESULTS_TABLE.csv](#)

CORE ARCHITECTURES & EMBEDDED SPECIFICATIONS:

- MultiScale-TelemetryAE (895p): Parallel Conv1D $k=[3,7,11]$, Enc2 12->6, Dec1 6->12, Dec2 12->1
- Distilled Micro-Student (421p): Single-branch Conv1D 1->8->4->8->1 (Dark Knowledge Distillation from MAML Teacher)
- Flash Storage: 3.50 KB FP32 (895p) / 1.64 KB FP32 (421p)
- Quantized SRAM: 895 Bytes INT8 (895p) / 421 Bytes INT8 (421p) [Fits in 0.22%-0.45% of 192 KB STM32F4 SRAM]
- Inference Latency: 1.2 ms / window on ARM Cortex-M4 @ 168 MHz (<0.15% CPU load at 1 Hz downlink)

DATA PIPELINE & ZERO-LEAKAGE CALIBRATION:

- Ingestion Datasets: NASA SMAP/MSL (81 ch), SKAB (32 series), SMD (28 traces), ESA OPS-SAT (8 ch), ESA-ADB (20k slice)
- Preprocessing: RobustScaler (Median / IQR) fitted strictly on Train 80% (Zero test-leakage guarantee)
- Windowing: $W = 100$ timesteps (100 s), Stride $S = 10$ (Train), Stride $S = 1$ (Inference point-by-point)
- Residual Scoring: $e_t = (x_t - \hat{x}_t)^2$ with Exponential Weighted Moving Average smoothing (span = 10)
- Threshold Calibration: Out-of-sample empirical quantile cutoff $\tau^* = \text{Quantile}(e_{\text{val}}, 0.985)$ on Val 20%

KEY VERIFIED BENCHMARKS (STRICT POINT-WISE RAW-F1):

- NASA SMAP/MSL (81 channels): MultiScale-895p = 0.3455 (5-seed: 0.3421 +/- 0.0084), PR-AUC = 0.4850
- NASA FFT-Augmented MultiScale: Raw-F1 = 0.3561 +/- 0.0076, Aff-F1 = 0.5340, PA-F1 = 0.8790
- Distilled Student (10-Shot Adapt): Raw-F1 = 0.3102, PA-F1 = 0.8643 (Segment Hit Rate: 80/104 = 76.92%)
- SKAB Testbed (32 Series): Multivariate 2,911p = 0.8141 +/- 0.0034 vs Univariate = 0.6869 +/- 0.0089 (+18.54%, $p=0.00001$)
- Baseline Comparison: Outperforms PatchTST (53.3k, 0.1523 Raw-F1, $p < 10^{-6}$) and USAD (27.8k, 0.1714, $p < 10^{-5}$)

VIVA DEFENSE CHEATSHEET & RAPID ANSWERS:

- Why not Transformers? Attention matrices scale $O(W^2)$ causing immediate OOM on sub-256 KB microcontrollers.
- Why is Raw-F1 ~0.35 when papers report 0.90+? Literature uses Point Adjustment (PA-F1), which awards full segment credit for single-point alarms. Under PA-F1, our model achieves 0.8643. True point-wise Raw-F1 is 0.14-0.35.
- Why parallel kernels $k=[3,7,11]$? Captures fast electrical/SEU spikes ($k=3$) and slow thermal drift ($k=11$) in parallel.
- Why not Gaussian 3-sigma? Telemetry residuals are heavy-tailed and multimodal; 3-sigma over-elevates thresholds (0.0537 F1).